



Research Article

Ant cuticle image classification using texture analysis: a comparative study

Noah Gardner¹, John Paul Hellenbrand², Anthony Phan¹, Haige Zhu¹, Zhiling Long⁴, Min Wang³, Clint Penick², and Chih-Cheng Hung^{1*}

¹ Laboratory of Machine Vision and Security Research, College of Computing and Software Engineering, Kennesaw State University, Marietta GA, USA

² College of Science and Mathematics, Kennesaw State University, Kennesaw, GA, USA

³ Department of Mathematics, Kennesaw State University, Kennesaw, GA, USA

⁴ Department of Mathematics, Science, and Informatics, Mercer University, Atlanta, GA, USA

* Correspondence: chung1@kennesaw.edu

Academic Editor: Pasi Fränti

Abstract: Ant cuticle texture presumably provides some type of function, and therefore is useful to research for ecological applications and bioinspired designs. In this study, we employ statistical image texture analysis and deep machine learning methods to automatically group similar ant species based on morphological features. We establish a public database of ant cuticle images for research. We provide a comparative study of the performance of image texture classification and deep machine learning methods on ant cuticle images. Our results show that the deep learning methods are giving higher accuracy in recognizing ant cuticle textures. Our experiments also reveal that the deep learning networks designed for image texture performs better than the general deep learning networks.

Keywords: texture analysis; image processing; classification; machine learning; ant cuticle images; ecology

1. Introduction

Insects compose half of biodiversity and rank among the most dominant organisms in terrestrial ecosystems [40]. A key factor for the ecological success of insects is their exoskeleton, also known as cuticle. The cuticle protects insects from predation, provides structural support, prevents desiccation, and serves as a canvas for advertising visual and chemical signals [15]. Research has focused on the macrostructures and internal chemical components that describe the functionality of the exoskeleton [16]. More recent work is being done to understand the functional aspects of external cuticle micro sculpturing [36,43].

Due to the extensive number of insect species, manual exploration of insect-based information is difficult and often requires expert training. Therefore, automated entomology has attracted biologists and computer scientists, and is expected to be a major contribution to the future of insect-based research [35]. One of the most commonly used data types for insect analysis is image data. To develop an image-based system for insect analysis, we can take advantage of existing work in general image analysis methods.

We examine ants (*Formicidae*) as they display an extreme diversity of cuticle micro sculpturing across all subfamilies. Microsculpture ranges from parallel longitudinal ridges to deep oval impressions to erratic protuberances. Microsculpture has arisen convergently and independently throughout ant's evolutionary history, which suggests that these are complex traits undergoing selection. Cuticle microsculpture on ants may help increase strength and rigidness, resist abrasion, increase internal and external surface area, resist microbial growth, and rear beneficial anti-biotic producing bacteria [4,8,25]. These specific functions may be associated with certain sculpturing types. In order to analyze those functions, it is necessary to segment the image and group similar textures for further analysis.

There is a large variety of ant species and most species are diverse in terms of size, shape, behaviors, and cuticle textures [references]. Based on our initial observation of the cuticle sculpturing in ants, it appears that some type of textural structures can be identified in images. Therefore, this image dataset may be worth to be explored from the image texture analysis point of view for the further knowledge discovery. Image texture analysis has been used in image processing for many decades [references]. Due to its diversity of images on scales, smoothness, coarseness, microtexture, and macrotecture occurring in images, it is a difficult problem for image textural interpreters. With a rapid development of deep machine learning, deep neural networks have been frequently used in image texture analysis [references].

Our contribution is to establish an ant image dataset which can be used to explore the categorization of different ant cuticle textures using image texture analysis methods. (comment: we may need to say how the dataset is established) Our experiment shows that the deep learning networks with image texture embedded in the design give a better accuracy than the K-views clustering method.

This paper is organized as below: Section 2 reviews the sculpturing identification in ants, and image texture analysis and insect classification. Section 3 describes the dataset preparation and image texture analysis methods used in our experiments. Section 4 gives the experimental results, visualization of both raw and processed images and analysis. The conclusion then follows.

2. Related work on sculpturing identification in entomology, image texture analysis, and insect classification

Taxonomists have developed extensive terminology describing ant cuticle microsculpture as it is often a useful diagnostic trait to distinguish between closely related ant species [2, 11]. The definitive text on ant cuticle terminology *The Glossary of Surface Sculpturing* contains over 100 terms to describe the cuticle sculpturing patterns of ants [19]. Despite the numerous ways to describe microsculpture, in general, microsculpture diversity falls into two classes - smooth and rough. Smooth ants are free of any obvious microsculpture and may have small hairs that give the cuticle a dull

appearance. Rough ants have a variety of raised lines or depressed textures on the cuticle. The entirety of ant microsculpture is represented within these two classes.

Cuticle sculpturing in ants has been explored thoroughly from a taxonomic perspective; however, the function of nano and microstructures on insect exoskeletons is a developing topic in entomology. Watson et al. review the literature of cuticle nano and microstructure function and propose 21 possible functions associated with these structures [43]. Many of these functions are related to structures not found in ants such as scales and nanostructures. The review does include functions that may relate to ants such as friction control, enhanced surface area, and increased hardness. Watson et al. also describe seven types of cuticle structures ranging from hairs and scales to nano and micro structures [43]. The cuticle sculpturing of ants seems to fall within one type - complex microstructures.

The five broad functional groupings developed that describe the complex microstructures found on ants were derived from reviewing the variation of cuticle sculpturing across ants. These functional groupings were also reflected in Harris as many terms could be grouped together based on similar definitions and comparing the scanning electron microscope photographs provided in the publication [19].

There is often substantial overlap among terms, and closely related cuticle patterns are likely to be functionally similar. For example, the definition for imbricate is, partly overlapping and appearing like shingles on a roof or scales on a fish, which is difficult to distinguish from the definition of tessellate, made up of squares like a chess board, either in sculpturing or in color. However, we note that the terms used to describe the cuticle patterns are similar to the terms used to describe textural patterns in texture analysis.

Image texture analysis is an important part of image interpretation as many objects appearing in the textural patterns either partially or completely inside an image. This phenomenon appears frequently in many images we encountered in our environment as almost everybody is using the iPhone to capture the photos. Scientists are also studying their subjects depending on taking the images remotely such as satellite images. With an abundant number of textural images surrounding us, it is essential to have a reliable automatic tool for the textural image interpretation.

Computer methods are developed for textural images based on the spatial relationship of gray levels of pixels encoded in the texture [18, 23]. Those methods by de-facto follow the steps used in the statistical pattern recognition [12]. Therefore, textural features are extracted, and the classification and segmentation methods are used for the interpretation either using the supervised or unsupervised approaches. The interpretation method is usually taken and modified from the traditional pattern recognition methods. For example, the gray-level co-occurrence matrix (GLCM) and local binary patterns (LBP) are two of methods for extracting features [17, 37].

Similar to the traditional pattern recognition, classical image texture analysis methods can be grouped into four categories: statistical, structural, model-based, and transform-based methods [1]. Among all of the categories, the statistical method is often used by employing the method from the pattern recognition principle. In doing so, an image texture is represented as feature vector using feature extraction methods and then fed into the algorithm. Markov random field models are also studied for textural image interpretation [7, 20]. The transform-based methods use some functions to decompose an image texture into a set of basic feature images. Gabor filters and Wavelet expansions are two of the widely used approaches [3].

While most feature vectors are established based on a single pixel, Liu et al. gave a comprehensive

survey on textural characterization in which they call this feature extraction as texture representation [30]. Their survey on Bag of Words and Convolutional Neural Networks (CNNs) [27] are the spatial relationship concepts. The K-views was developed for taking the spatial information in the clustering approach [23].

In the past decade, deep learning using CNNs has emerged as the mainstream technology for image analysis. Following this trend, various CNN-based network architectures have been designed to specifically characterize textural images. Among these, the Fisher-vector CNN descriptor (FV-CNN) proposed by Cimpoi et al. [6] is widely accepted as one of the most important pioneering works. It applies FV pooling to deep features obtained via a CNN pre-trained using the ImageNet [27] to obtain encoded features for texture classification. Although it is capable of achieving much improved classification accuracy comparing to traditional hand-crafted texture features, FV-CNN does not support an end-to-end learning, where feature extraction, encoding, and classification are separated from each other. To achieve an end-to-end learning, Zhang et al. [48] proposed the deep texture encoding network (DeepTEN), in which a novel texture encoding layer is added to a standard CNN architecture. Then, Xue et al. [44] constructed the deep encoding pooling network (DEP), which improves over DeepTEN by integrating local spatial characteristics into the texture representation. Based on DeepTEN and DEP, Hu et al. [22] further developed the multi-level texture encoding and representation (MuLTER) network, which embeds a learnable encoding module at each convolutional layer so that encoding is performed for both low-level and high-level features, yielding a multi-level texture representation.

Other network architectures for end-to-end texture learning are also available. For example, in the deep multiple-attribute-perceived network (MAP-Net) [47], multiple perceptual attributes are progressively learned in a mutually reinforced manner through multiple branches. In the deep structure-revealed network (DSR-Net) [46], inherent structural representation for a texture pattern is obtained by employing a primitive capturing module to learn spatial primitives and a dependency learning module to capture the dependency among the primitives. In [34], a residual pooling layer consisting of a residual encoding module and an aggregation module is used to generate discriminative features of low dimensions. In [38], a histogram layer is designed to compute local spatial distribution of CNN features. In [5], an innovative aggregation module is presented to exploit statistical self-similarity across layers. All these architectures customize the standard CNN structure to accomplish the characterization of certain spatial, visual, and statistical traits unique to textural images.

Compared with the traditional method in which a kernel must be designed by an engineer for extracting features, the deep learning networks can automatically extract the features through the training. In addition, the deep learning networks achieve the higher accuracy than that of the classical approaches, although the deep networks require much more data for training.

Proposed insect classification methods seek to classify insects at different hierarchical levels, such as species, genus, family, and order. Additionally, some methods may classify insects at a combination of different hierarchical levels. Insect classification methods can be applied to a variety of fields. In agriculture, insect classification methods can be used to identify the presence of pest insects in crops, which can inform crop managers in their choice of pesticides and help prevent crop loss [26, 31].

Feng et al. [10] apply an automated system to classify moth images based on semantic related visual attributes, which are defined as a pattern on the moth wings. Feng et al. [10] use a custom texture descriptor based on the combination of GLCM and *scale-invariant feature transform* (SIFT)

features [14, 33]. The method proposed by Feng et al. [10] is used to classify 50 different moth species across 8 families [10]. The results from Feng et al. [10] suggest that traditional feature extraction techniques for the semantic visual attributes of the moth wings are sufficient for training a classifier to classify an image between 10 randomly selected moth species.

Urteaga et al. [42] use machine learning methods in order to classify images between two different scorpion species: *Centruroides limpidus* and *Centruroides noxius*. After applying background distinction based on dynamic color threshold, Urteaga et al. [42] apply feature extraction to extract features from the separated scorpion image such as aspect ratio, rectangularity, and compactness. Urteaga et al. [42] apply three different classification models to classify the image as one of the species: Artificial Neural Network, Regression Tree, and Random Forest classifiers [42]. The results from Urteaga et al. [42] show that after background removal, characteristics from the entire body of the scorpion can be used to create a binary classifier that can classify the image as one of the two species.

Lim et al. [29] apply a CNN-based algorithm for insect classification. Lim et al. [29] classify a subset of insect species and families based on the classes available in the ImageNet dataset. ImageNet is a widely used dataset of images labeled by experts with millions of images and thousands of categories [9]. In the ImageNet dataset, there are some categories that specify the class of the insect on a species level, e.g. *monarch butterfly* and *ringlet butterfly* as well as some categories that specify the class of the insect on a family level, e.g. *ant*, *fly*, and *bee* [9]. Lim et al. [29] use a modified AlexNet architecture and experiment with different numbers of kernels and their effect the performance of the model. Glick et al. [13] employ a similar approach by classifying 277 insect classes from ImageNet using a hierarchical convolutional neural network. The results from Lim et al. [29] and Glick et al. [13] suggest that a CNN is capable of differentiating between different hierarchical classes of insects.

3. Ant cuticle image dataset

In this section, we describe the creation of the custom dataset used in this research. In our dataset, we use ant head images from AntWeb [39] and define two categories for them based on the appearance of the cuticle texture: *rough* and *smooth*. We also refer to these ant head images, which describe the pose of the ant, as ant cuticle images, due to the clear cuticle texture present on the head. Some randomly selected images from each category are shown in Figures 1 and 2.

To begin, a master spreadsheet was created with the 2,499 different ant species to be identified for the primary dataset. The team was trained to identify cuticle sculpturing through a process which consisted of one 45-minute introductory lesson explaining the project and texture categories. Then, the team was given a training set of photos to identify from the genus *Polyrhachis*. The sculpture identification protocol describes the two primary categories: rough and smooth.

Initially, the sculpture identification protocol had 5 subcategories of cuticle texture: smooth, punctate, striate, reticulate, and tuberous. For simplicity, we work only with the two main categories. The training set identifications were reviewed together as a group by the assistants. Once training was complete, assistants were assigned the same genera of ants to identify independently each week. A weekly meeting was held to discuss identifications and assign new ones. These identifications were collected in the master spreadsheet and the identifications were assigned to individual ant species on a majority basis.

To collect the images, the assistants followed the taxonomy information available in the master



Figure 1. Examples of smooth cuticle texture ant images in the dataset after center cropping, from AntWeb [39].



Figure 2. Examples of rough cuticle texture ant images in the dataset after center cropping, from AntWeb [39].

spreadsheet to the appropriate AntWeb* page. In many cases, there are multiple ant head images of the same species, and occasionally there are multiple image resolution available from a single image. To simplify the data collection process, the assistants were instructed to download the first ant head image of the species being identified in the highest resolution possible. Each image was named with an identifier that corresponds with the row number in the master spreadsheet. The same ant head images that were downloaded in the data collection phase were the same ones used in the sculpture identification protocol. Ant species which did not have any images of the head were excluded from the dataset. Additionally, ant species which only had a head image of a queen ant were excluded from the dataset.

Ant specimen images taken from AntWeb [39] are created by different photographers and therefore have different attributes, such as environment, resolution, and lighting. In the ant head images, the ant head is in the center of the image and the body is pointing away from the camera. The focus of the ant head image is centered on the head, with the background and image artifacts from the ant body typically blurred. In most ant head images, there is a bar which indicates the scale of the image due to the variety in the sizes of different ant species. In a few ant head images, there exists some text denoting the specimen identifier and other information. In terms of texture, some ant specimens are very old, so their head images have other abnormalities such as cracks in the cuticle and the presence of dust.

The completed dataset of contains 2,499 images. 1072 samples of rough textured ant cuticle textures comprise 43% of the dataset. The remaining 1427 samples of smooth textured ant cuticle textures comprise 57% of the dataset. Due to the variety of the ant head image attributes, we apply simple preprocessing before the images are used in our model. We want the images to have a uniform size for simplicity in our classification process. Since the ant head images are typically centered in the image, we apply a center crop to each image to create a square image of the same size. Once the image is square, we resize each image to a fixed size of 256x256 pixels. We leave other discrepancies in the images untouched. A summary of the resulting dataset is shared in Table 1.

Table 1. Summary of the dataset.

Total images	2499
Rough images	1072
Smooth images	1427
Image size	256x256

4. Methods for comparative analysis and experimental setup

The K-means algorithm is one of widely used clustering algorithms in the pattern recognition [32]. It is a single pixel based classification algorithm for images. The K-views is an algorithm that uses several characteristic views for the classification of images [24]. The K-views algorithm is suitable for classifying image textures that have basic local patterns repeated in a periodic manner. In contrast to the K-means, K-views looks at the neighboring pixels for providing the spatial relationship for classification. Several variations of the basic K-views model have been proposed [28, 45].

*<https://www.antweb.org/>

For the feasibility study in this ant image datasets, we use both K-means and K-views in our experiments as both are statistical clustering methods. The experiments may indicate how well the statistical approach will do in this type of textural images. In pre-processing the dataset, we apply the Gabor filter to transform the gray-scale ant image into four Gabor-filtered channels, subtract the gray-scale image from each channel, take the absolute value of the resulting images, and then normalize them. This pre-processing will highlight some of characteristics features such as edge-like in the image. This step is very useful in feeding the datasets to the K-views algorithms. A 10-channel Gabor filter is applied in the pre-processing step for the K-means algorithm. The theta θ parameters in the Gabor filters used are multiples of $\pi/10$ radians starting from 0 to π radians. The Gabor filter lambda λ parameter is equal to π .

Our first deep learning model is *visual geometry group* (VGG), a convolutional neural network that takes advantage of very small convolutional filters in a deep network architecture [41]. We compare four architectures of VGG: VGG11, VGG13, VGG16, and VGG19. The primary difference between the architectures is the number of layers in each model. Our second model is *residual network* (ResNet), a deep network architecture that includes shortcut connections between layers (residual connections) [21]. We compare three architectures of ResNet: ResNet18, ResNet50, and ResNet101. Again, the primary difference between the architectures is the number of layers in each model.

For our ResNet models, we have two versions: randomized and pretrained. The randomized version is the same architecture, but the weights are randomly initialized. The pretrained version has weights from training on the CIFAR dataset, an image dataset with 1000 classes. In this case, we are fine-tuning the pretrained model. For VGG, we are only using the randomized version. The base VGG architecture also has an output layer of size 1000. Since we are working with a binary classification problem, we modify the architecture for all models to have an output layer of size 2. Each model is trained over 100 epochs, using stochastic gradient descent with momentum. The batch size is set to 16 images. We apply a learning rate of 0.001 and momentum parameter of 0.9.

In addition to the classical texture analysis methods and the deep learning models for general image analysis, we also examine deep learning analysis methods specifically designed for textural images. For this examination, we adopt the deep residual pooling (DRP) network developed by Mao et al. [34]. The DRP framework consists of a unique residual pooling layer, which is formed by a residual encoding module followed by an aggregation module. The residual encoding module extracts relevant spatial information, while the aggregation module performs averaging to obtain orderless low-dimension features for classification.

According to [34], the DRP network can be used either with a single residual pooling layer, applied right before the classifier layer, or with multiple residual pooling layers, applied to the output of multiple convolutional layers before feeding the combined pooling results into the classifier layer. Additionally, an auxiliary classifier layer may also be used. In our experiments, we follow these strategies. For each scenario, we conduct two experiments. One experiment employs weights that are randomly initialized, and the other uses weights that are fine tuned over those from a pretrained model.

To handle the imbalance of the dataset, we apply undersampling for each class for the training dataset. By using random stratified sampling, we construct a training set with 800 images per class. The remaining images are randomly split between test and validation, which turns out to roughly a 60%/20%/20% train, test, and validation data split. With 272 rough samples and 627 smooth samples left over after the stratified split, the test dataset has roughly 136 rough samples and 313 smooth

samples. Since these leftover samples are split with code by 50% there will be some rounding variance and therefore the test dataset built at run-time will not always have exactly the same number of samples.

We evaluate the performance of the models according to standard evaluation methods. Since we are working with a binary classification problem, we use a standard confusion matrix to evaluate the accuracy, precision, and F1 score. We also apply Grad-CAM with manual inspection to visualize the activation weights for classified images to visualize which features lead to the classification result. Finally, we apply t-SNE to visualize the separation learned for the model to further analyze the classifications made by the model.

Experiments are run on an Ubuntu 18.04 LTS Lambda Labs GPU server. The server contains 8 NVIDIA GeForce RTX 2080 Ti graphics cards with 12GB of memory each. The server uses an Intel Xeon Silver 4116 with 48 total threads and maximum frequency of 3.000 GHz, and has 256GB of RAM.

5. Results and analysis

To begin, we share the results on the each algorithm specified in the methodology section using the evaluation metrics described. The results are shown in Table 2. It should be noted that due to the class imbalance in the dataset, the F1 score is the preferable metric to the accuracy.

Table 2. Experimental results of different algorithms

Model	Accuracy	Precision	Recall	F1
K-Means	0.47	0.07	0.35	0.11
K-Views (15x15)	0.56	0.58	0.56	0.57
K-Views (17x17)	0.62	0.79	0.59	0.68
K-Views (19x19)	0.62	0.80	0.59	0.68
K-Views (25x25)	0.52	0.71	0.51	0.59
ResNet18	0.80	0.72	0.67	0.69
ResNet18 (fine-tuned)	0.88	0.87	0.77	0.82
ResNet50	0.78	0.59	0.67	0.62
ResNet50 (fine-tuned)	0.87	0.83	0.78	0.80
ResNet101	0.77	0.63	0.62	0.62
ResNet101 (fine-tuned)	0.89	0.87	0.78	0.82
VGG11	0.80	0.73	0.66	0.69
VGG13	0.82	0.69	0.71	0.69
VGG16	0.80	0.65	0.67	0.65
VGG19	0.81	0.67	0.68	0.64
DRP Multi-layer	0.87	0.89	0.92	0.91
DRP Multi-layer (fine-tuned)	0.88	0.89	0.93	0.91
DRP Single-layer	0.88	0.90	0.94	0.92
DRP Single-layer (fine-tuned)	0.88	0.90	0.93	0.91
DRP Single-layer Auxiliary	0.87	0.89	0.92	0.90
DRP Single-layer Auxiliary (fine-tuned)	0.89	0.90	0.94	0.92

The results in the previous sections show that the fine-tuned ResNet models outperform the VGG and randomly initialized ResNet models on the task of ant head image classification. Additionally, the DRP models perform better than the general deep learning models VGG and ResNet on the task of ant head image classification. The fine-tuned DRP single-layer model with auxiliary classifier performed the best with an average F1 score of 0.92. We further analyze the separation learned by both ResNet101 models in the following section.

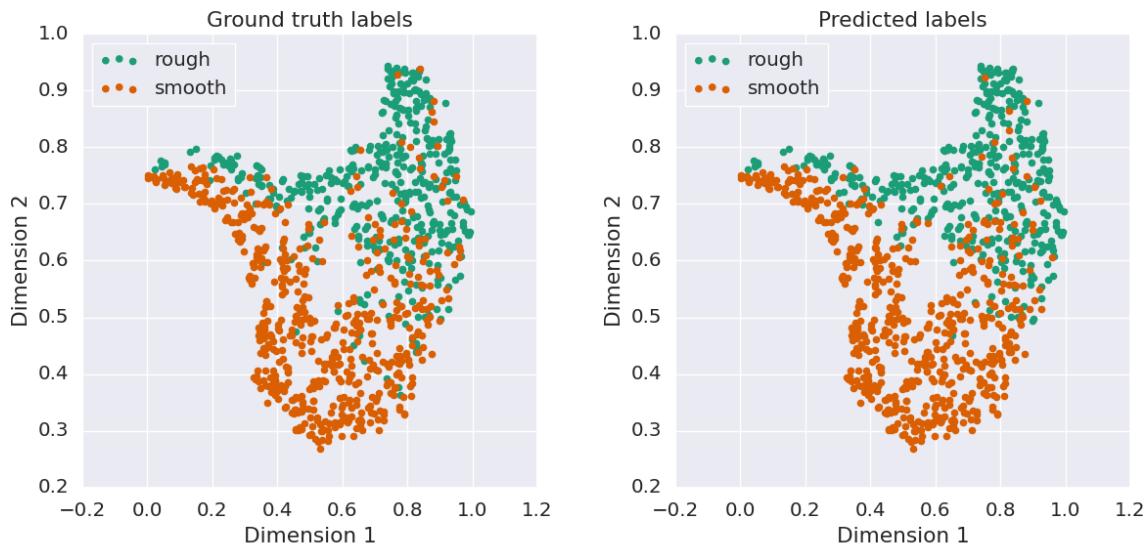


Figure 3. t-SNE visualization of the embeddings of the second to last layer of the randomly initialized ResNet101 model trained on ant head image dataset.

In this section, we provide visualization of the fine-tuned ResNet101 model and the randomly initialized ResNet101 model using t-SNE dimensionality reduction. First, we run the dataset preprocessing method and initialize both models. Then, both models are trained according to the training parameters and the state of each model is saved. To visualize the deep extracted features, we modify each model to obtain the embeddings of the second to last layer. Then, we use the t-SNE algorithm to reduce the dimensionality of the embeddings to 2 dimensions. We plot side-by-side the ground truth and predicted labels for each model. Figure 3 shows the results of the trained randomly initialized model and Figure 4 shows the results of the fine-tuned model. Based on the visual results of the t-SNE visualization, we can see that the fine-tuned model learned a stronger separation of the two classes, which reinforces the results that the fine-tuned model received a higher average accuracy.

Next, we provide some visual analysis of some correctly and incorrectly classified images using GradCAM. We provide two categories and two subcategories in our analysis. The two categories are correct and incorrect classification. Regardless of the feature activation map, the correctly classified images have the same predicted label as the ground truth, and incorrectly classified images have different predicted labels. The two subcategories are expected and non-expected feature activation. In the expected case, the features that are used to compute the classification are the same as the features used by the assistants in the sculpture identification process. In general, the features used by the assistants are the textures of the cuticle on the ant head. In the non-expected case, the features used to compute the classification are not from the head, for example, from the background, extraneous text, or

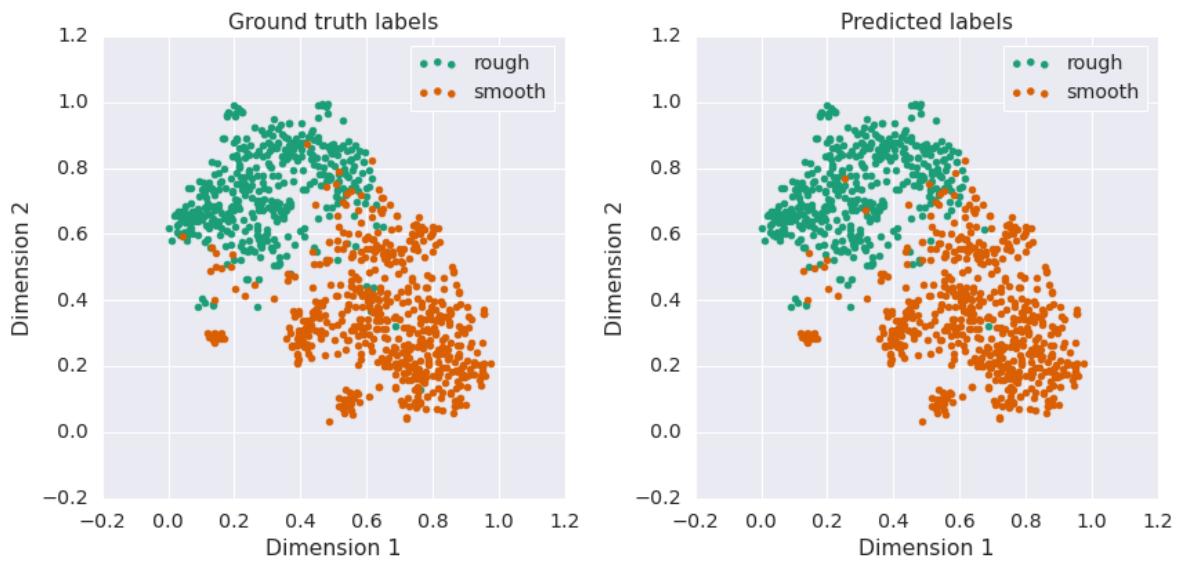


Figure 4. t-SNE visualization of the embeddings of the second to last layer of the fine-tuned ResNet101 model trained on ant head image dataset.

the body of the ant. We used randomly selected images from the dataset and the fine-tuned ResNet101 model to perform the analysis. We show the GradCAM results in Figures 5 and 6. The left image shows the preprocessed image input to the model. The right image shows the GradCAM output based on the classification. Four specimens were selected randomly from each category and subcategory.

Correctly classified images which use the expected features show the ideal performance of the model. Incorrectly classified images which use the expected features should be further analyzed. In essence, the model in this situation knows *where* to look, but not *what* to look for. In Figure 6, there are results where the features activated are mostly in the correct location on the ant head, and the rough texture is clearly visible, yet the model predicts the incorrect class *smooth*. Similarly in Figure 6, there are results where the features activated are also mostly in the correct location, yet the model predicts the incorrect class *smooth*. In this case, it may be due to the pose of the ant being slightly different from the average pose. In the incorrectly classified images with non-expected features, analysis shows that the model is unable to find *where* to look, and obtains feature information from other parts of the ant or the background. Cases where the image was correctly classified using the non-expected features can basically be seen as noise. In order to further analyze this class, we should introduce some parameter such as model confidence to examine further.

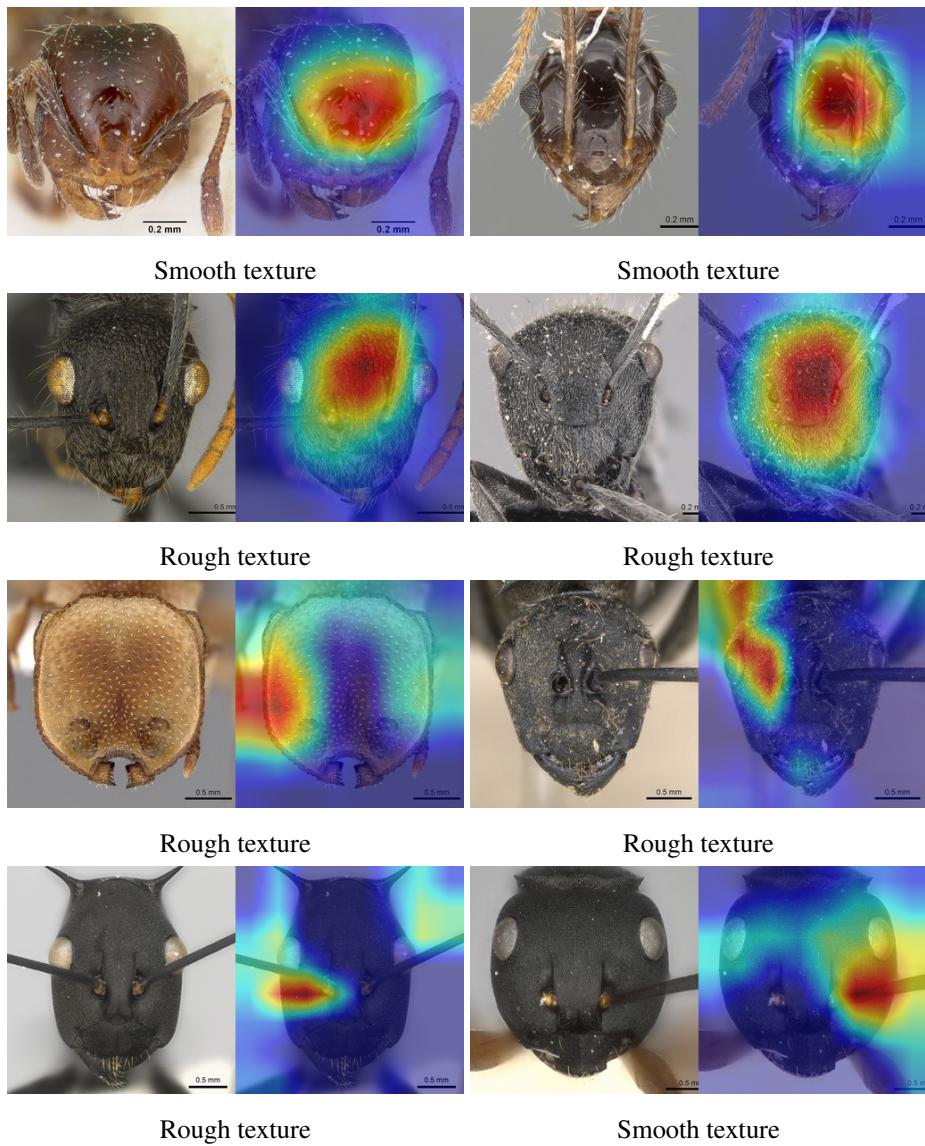


Figure 5. Correctly classified images.

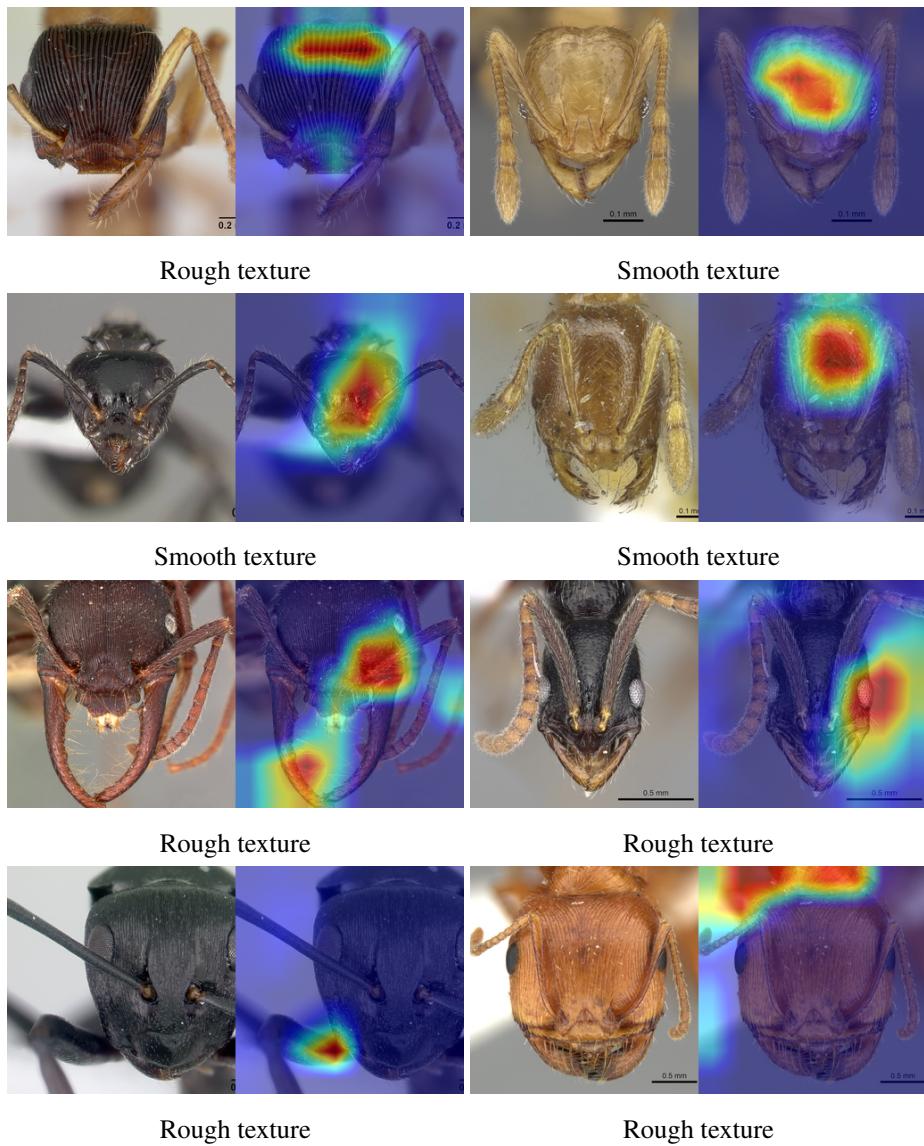


Figure 6. Incorrectly classified images.

6. Conclusion

Ant cuticle texture presumably has some function, but without the proper tools, evaluating the function based on thousands of species is infeasible. We have shown in this work that a deep learning approach and classical image texture analysis methods can be used to automatically categorize ants based on their cuticle texture, therefore supporting research on the evaluation of the function in future work. Our categorization system is novel in the field of automated insect identification due to the broad number of species captured by it. Additionally, a model that is pre-trained on a diverse image task such as ResNet can be transferred to our domain of texture analysis. However, a deep learning algorithm created specifically for the domain of texture analysis had the best results in our experiments. In future work, we will continue to explore texture classes present in the dataset which are not captured by the binary class system. All code is publicly available on GitHub (<https://github.com/ngngardner/cuticulus>).

References

1. M. H. Bharati, J. J. Liu and J. F. MacGregor, Image texture analysis: methods and comparisons, *Chemometrics and intelligent laboratory systems*, **72** (2004), 57–71.
2. B. B. Blaimer, Taxonomy and Natural History of the Crematogaster (Decacrema)-group (Hymenoptera: Formicidae) in Madagascar, *Zootaxa*, **2714** (2019), 1, URL <https://biotaxa.org/Zootaxa/article/view/zootaxa.2714.1.1>.
3. A. C. Bovik, M. Clark and W. S. Geisler, Multichannel texture analysis using localized spatial filters, *IEEE transactions on pattern analysis and machine intelligence*, **12** (1990), 55–73.
4. A. Brückner, M. Heethoff and N. Blüthgen, The relationship between epicuticular long-chained hydrocarbons and surface area - volume ratios in insects (Diptera, Hymenoptera, Lepidoptera), *PLOS ONE*, **12** (2017), e0175001, URL <https://dx.plos.org/10.1371/journal.pone.0175001>.
5. Z. Chen, F. Li, Y. Quan, Y. Xu and H. Ji, Deep Texture Recognition via Exploiting Cross-Layer Statistical Self-Similarity, in *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, Nashville, TN, USA, 2021, 5227–5236, URL <https://ieeexplore.ieee.org/document/9577662/>.
6. M. Cimpoi, S. Maji and A. Vedaldi, Deep filter banks for texture recognition and segmentation, in *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, Boston, MA, USA, 2015, 3828–3836, URL [http://ieeexplore.ieee.org/document/7299007/](https://ieeexplore.ieee.org/document/7299007/).
7. G. R. Cross and A. K. Jain, Markov random field texture models, *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 25–39.
8. C. R. Currie, Coevolved Crypts and Exocrine Glands Support Mutualistic Bacteria in Fungus-Growing Ants, *Science*, **311** (2006), 81–83, URL <https://www.sciencemag.org/lookup/doi/10.1126/science.1119744>.
9. J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li and L. Fei-Fei, ImageNet: A large-scale hierarchical image database, in *2009 IEEE Conference on Computer Vision and Pattern Recognition*, 2009, 248–255, ISSN: 1063-6919.

-
10. L. Feng and B. Bhanu, Automated identification and retrieval of moth images with semantically related visual attributes on the wings, in *2013 IEEE International Conference on Image Processing*, 2013, 2577–2581, ISSN: 2381-8549.
 11. B. L. Fisher and S. P. Cover, *Ants of North America: a guide to the genera*, University of California Press, Berkeley, 2007, OCLC: ocm80180487.
 12. K. Fukunaga, *Introduction to statistical pattern recognition*, Elsevier, 2013.
 13. J. A. Glick and K. Miller, Insect Classification With Heirarchical Deep Convolutional Neural Networks Convolutional Neural Networks for Visual Recognition, 2016.
 14. C. C. Gotlieb and H. E. Kreyszig, Texture descriptors based on co-occurrence matrices, *Computer Vision, Graphics, and Image Processing*, **51** (1990), 70–86, URL <https://linkinghub.elsevier.com/retrieve/pii/S0734189X05800635>.
 15. P. J. Gullan and P. S. Cranston, *Insects: an Outline of Entomology.*, Wiley, Hoboken, 2009, URL <http://qut.eblib.com.au/patron/FullRecord.aspx?p=233169>, OCLC: 1048427241.
 16. S. Gunderson and R. Schiavone, The insect exoskeleton: A natural structural composite, *JOM*, **41** (1989), 60–63, URL <http://link.springer.com/10.1007/BF03220386>.
 17. R. M. Haralick, K. Shanmugam and I. Dinstein, Textural Features for Image Classification, *IEEE Transactions on Systems, Man, and Cybernetics*, **SMC-3** (1973), 610–621, URL <http://ieeexplore.ieee.org/document/4309314/>.
 18. R. M. Haralick and L. G. Shapiro, Image segmentation techniques, *Computer vision, graphics, and image processing*, **29** (1985), 100–132.
 19. R. A. Harris, A Glossary Of Surface Sculpturing., *Occasional Papers in Entomology*, URL <https://zenodo.org/record/26215>.
 20. M. Hassner and J. Sklansky, The use of markov random fields as models of texture, in *Image Modeling*, Elsevier, 1981, 185–198.
 21. K. He, X. Zhang, S. Ren and J. Sun, Deep Residual Learning for Image Recognition, in *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016, 770–778, ISSN: 1063-6919.
 22. Y. Hu, Z. Long and G. AlRegib, Multi-Level Texture Encoding and Representation (Multer) Based on Deep Neural Networks, in *2019 IEEE International Conference on Image Processing (ICIP)*, IEEE, Taipei, Taiwan, 2019, 4410–4414, URL <https://ieeexplore.ieee.org/document/8803640/>.
 23. C.-C. Hung, E. Song and Y. Lan, *Image Texture Analysis: Foundations, Models and Algorithms*, Springer International Publishing, Cham, 2019, URL <http://link.springer.com/10.1007/978-3-030-13773-1>.
 24. C. Hung, S. Yang and C. M. Laymon, Use of characteristic views in image classification, in *16th International Conference on Pattern Recognition, ICPR 2002, Quebec, Canada, August 11–15, 2002*, IEEE Computer Society, 2002, 949–952, URL <https://doi.org/10.1109/ICPR.2002.1048462>.
 25. R. A. Johnson, A. Kaiser, M. Quinlan and W. Sharp, Effect of cuticular abrasion and recovery on water loss rates in queens of the desert harvester ant *Messor pergandei*, *Journal of Experimental Applied Computing and Intelligence*

-
- Biology*, **214** (2011), 3495–3506, URL <https://journals.biologists.com/jeb/article-214/20/3495/10482/Effect-of-cuticular-abrasion-and-recovery-on-water>.
26. T. Kasinathan and S. R. Uyyala, Machine learning ensemble with image processing for pest identification and classification in field crops, *Neural Computing and Applications*, **33** (2021), 7491–7504, URL <https://doi.org/10.1007/s00521-020-05497-z>.
 27. A. Krizhevsky, I. Sutskever and G. E. Hinton, ImageNet classification with deep convolutional neural networks, *Communications of the ACM*, **60** (2017), 84–90, URL <https://dl.acm.org/doi/10.1145/3065386>.
 28. Y. Lan, H. Liu, E. Song and C. Hung, An improved k-view algorithm for image texture classification using new characteristic views selection methods, in *Proceedings of the 2010 ACM Symposium on Applied Computing (SAC), Sierre, Switzerland, March 22-26, 2010* (eds. S. Y. Shin, S. Ossowski, M. Schumacher, M. J. Palakal and C. Hung), ACM, 2010, 959–963, URL <https://doi.org/10.1145/1774088.1774288>.
 29. S. Lim, S. Kim and D. Kim, Performance effect analysis for insect classification using convolutional neural network, in *2017 7th IEEE International Conference on Control System, Computing and Engineering (ICCSCE)*, 2017, 210–215.
 30. L. Liu, J. Chen, P. Fieguth, G. Zhao, R. Chellappa and M. Pietikinen, From BoW to CNN: Two Decades of Texture Representation for Texture Classification, *International Journal of Computer Vision*, **127** (2019), 74–109, URL <http://link.springer.com/10.1007/s11263-018-1125-z>.
 31. L. Liu, R. Wang, C. Xie, P. Yang, F. Wang, S. Sudirman and W. Liu, PestNet: An End-to-End Deep Learning Approach for Large-Scale Multi-Class Pest Detection and Classification, *IEEE Access*, **7** (2019), 45301–45312.
 32. S. Lloyd, Least squares quantization in pcm, *IEEE transactions on information theory*, **28** (1982), 129–137.
 33. D. G. Lowe, Distinctive Image Features from Scale-Invariant Keypoints, *International Journal of Computer Vision*, **60** (2004), 91–110, URL <http://link.springer.com/10.1023/B:VISI.0000029664.99615.94>.
 34. S. Mao, D. Rajan and L. T. Chia, Deep residual pooling network for texture recognition, *Pattern Recognition*, **112** (2021), 107817, URL <https://linkinghub.elsevier.com/retrieve/pii/S0031320321000042>.
 35. M. Martineau, D. Conte, R. Raveaux, I. Arnault, D. Munier and G. Venturini, A survey on image-based insect classification, *Pattern Recognit.*, **65** (2017), 273–284.
 36. S. Muthukrishnan, S. Mun, M. Y. Noh, E. R. Geisbrecht and Y. Arakane, Insect Cuticular Chitin Contributes to Form and Function, *Current Pharmaceutical Design*, **26** (2020), 3530–3545, URL <https://www.eurekaselect.com/182235/article>.
 37. T. Ojala, M. Pietikinen and T. Menp, Gray Scale and Rotation Invariant Texture Classification with Local Binary Patterns, in *Computer Vision - ECCV 2000* (eds. G. Goos, J. Hartmanis and J. van Leeuwen), vol. 1842, Springer Berlin Heidelberg, Berlin, Heidelberg, 2000, 404–420, URL http://link.springer.com/10.1007/3-540-45054-8_27.

38. J. Peeples, W. Xu and A. Zare, Histogram Layers for Texture Analysis, *IEEE Transactions on Artificial Intelligence*, 1–1, URL <https://ieeexplore.ieee.org/document/9652037/>.
39. V. Perrichot and B. Fisher, AntWeb: digitizing Recent and fossil insects for an online database of the ants of the world, in *Digital Fossil International Conference*, 2012, URL <https://hal-insu.archives-ouvertes.fr/insu-00805243>.
40. A. Sheikh, N. Rehman and R. Kumar, Diverse adaptations in insects: A review, *Journal of entomology and zoology studies*, **5** (2017), 343–350.
41. K. Simonyan and A. Zisserman, Very Deep Convolutional Networks for Large-Scale Image Recognition, *arXiv:1409.1556 [cs]*, URL <http://arxiv.org/abs/1409.1556>, ArXiv: 1409.1556.
42. J. C. Urteaga-Reyesvera and A. Possani-Espinosa, Scorpions: Classification of poisonous species using shape features, in *2016 International Conference on Electronics, Communications and Computers (CONIELECOMP)*, 2016, 125–129.
43. G. S. Watson, J. A. Watson and B. W. Cribb, Diversity of Cuticular Micro- and Nanostructures on Insects: Properties, Functions, and Potential Applications, *Annual Review of Entomology*, **62** (2017), 185–205, URL <http://www.annualreviews.org/doi/10.1146/annurev-ento-031616-035020>.
44. J. Xue, H. Zhang and K. Dana, Deep Texture Manifold for Ground Terrain Recognition, in *2018 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, IEEE, Salt Lake City, UT, USA, 2018, 558–567, URL <https://ieeexplore.ieee.org/document/8578163/>.
45. S. Yang and C. Hung, Image texture classification using datagrams and characteristic views, in *Proceedings of the 2003 ACM Symposium on Applied Computing (SAC), March 9-12, 2003, Melbourne, FL, USA* (eds. G. B. Lamont, H. Haddad, G. A. Papadopoulos and B. Panda), ACM, 2003, 22–26, URL <https://doi.org/10.1145/952532.952538>.
46. W. Zhai, Y. Cao, Z.-J. Zha, H. Xie and F. Wu, Deep Structure-Revealed Network for Texture Recognition, in *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, Seattle, WA, USA, 2020, 11007–11016, URL <https://ieeexplore.ieee.org/document/9157507/>.
47. W. Zhai, Y. Cao, J. Zhang and Z.-J. Zha, Deep Multiple-Attribute-Perceived Network for Real-World Texture Recognition, in *2019 IEEE/CVF International Conference on Computer Vision (ICCV)*, IEEE, Seoul, Korea (South), 2019, 3612–3621, URL <https://ieeexplore.ieee.org/document/9009995/>.
48. H. Zhang, J. Xue and K. Dana, Deep TEN: Texture Encoding Network, in *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, IEEE, Honolulu, HI, 2017, 2896–2905, URL <http://ieeexplore.ieee.org/document/8099792/>.